

Comp 6321 - Machine Learning - Assignment 4

Federico O'Reilly Regueiro

December 4th, 2016

Question 1: VC dimensions

1.a $[a, \infty)$

We can shatter a single point $p_1, p_1 \in \mathbb{R}$:

point	label	h
p_1	\oplus	$[a, \infty), a < p_1$
p_1	\ominus	$[a, \infty), a > p_1$

But if we have two points, $p_1, p_2 \mid p_1 < p_2, p_1 \in \oplus, p_2 \in \ominus$, then $[a, \infty)$ cannot shatter them. Therefore, for this class of hypothesis: $VC_{dim} = 1$

1.b $(-\infty, a]$ or $[a, \infty)$

Similarly to the previous question, we can shatter one point. Additionally, we can shatter two points, $p_0, p_1 \mid p_0 < p_1, p_0$:

point	label	h
p_1	\ominus	$(-\infty, a], a < p_1$
p_2	\ominus	
p_1	\ominus	$[a, \infty), p_1 < a < p_2$
p_2	\oplus	
p_1	\oplus	$(-\infty, a], p_1 < a < p_2$
p_2	\ominus	
p_1	\oplus	$[a, \infty), a < p_1$
p_2	\oplus	

However, three points $p_1, p_2, p_3, \mid p_1 < p_2 < p_3, p_1 \in \ominus, p_2 \in \oplus, p_3 \in \ominus$ cannot be shattered. Therefore, for this class of hypothesis: $VC_{dim} = 2$

1.c Finite unions of one-sided intervals

The union of more than one left-side interval $(-\infty, a] \cup (-\infty, b] \dots \cup (-\infty, n]$ is equivalent to a single left-side interval $(-\infty, \max(a, b, \dots n)]$. The same applies for one or more right-side intervals being equivalent to $[\min(a, b, \dots n), \infty)$. Therefore, this hypothesis class is of the form $(-\infty, a] \cup [b, \infty)$.

Since $\{(-\infty, a] \text{ or } [b, \infty)\} \subset \{(-\infty, a] \cup [b, \infty)\}$, we know this class of hypothesis to be capable of shattering 2 points. But once again, three points $p_1, p_2, p_3, \mid p_1 < p_2 < p_3, p_1 \in \ominus, p_2 \in \oplus, p_3 \in \ominus$ cannot be shattered with this class of hypothesis. Therefore, for this class: $VC_{dim} = 2$

1.d $[a, b] \cup [c, d]$

This class of hypothesis can shatter four points due to the following:

- a Any four positives can be correctly classified by a single interval as can any labeling with a single positive.
- b Any two positives and two negatives can be classified with two intervals, given that a single interval is assigned to each positive.
- c Labeling three positives and one negative will always yield at most two disjunct groups of contiguous positive labels, each of which can be contained in one of the two intervals.

However, if we have five points $p_1, p_2, p_3, p_4, p_5, | p_1 < p_2 < p_3 < p_4 < p_5, p_1 \in \oplus, p_2 \in \ominus, p_3 \in \oplus, p_4 \in \ominus, p_5 \in \oplus$ cannot be shattered with this class of hypothesis. Therefore, for this class: $VC_{dim} = 4$

1.e Unions of k intervals

Informally speaking, the maximum number of disjunct groups of positives in a set with p points, where p is even, is $\frac{p}{2}$. Thus k intervals would allow us to correctly label all positive points in a set with cardinality $2k$. We prove this is the case by induction:

Base step: One interval, $k = 1, h = [a, b]$, and two points, $p_1, p_2 | p_1 < p_2, p_1:$

point	label	h
p_1 p_2	\ominus \ominus	$[a, b], b < p_1$
p_1 p_2	\ominus \oplus	$[a, b], p_1 < a < p_2 < b$
p_1 p_2	\oplus \ominus	$[a, b], a < p_1 < b < p_2$
p_1 p_2	\oplus \oplus	$[a, b], a < p_1, p_2 < b$

We increase the set to three points with the following labels $p_1, p_2, p_3, | p_1 < p_2 < p_3, p_1 \in \oplus, p_2 \in \ominus, p_3 \in \oplus$ and note that it cannot be shattered. Therefore, for the base step $VC_{dim} = 2 = 2k$.

Now we suppose that for the union of k intervals, VC dimension is $2k^1$, then we need to prove that with $k + 1$ intervals we are able to shatter $2(k + 1)$.

Firstly we note that the most *difficult* configuration to classify would be an alternation of \oplus and \ominus points, since it would require using each one of the k intervals to classify a single point each; any other configuration would require less than k intervals and we would have some *leftover* intervals to be consumed in classifying newly inserted points.

Inductive step: We add points p_{2k+1}, p_{2k+2} , with no inequality constraints, to the $2k$ points shattered with k intervals. Without loss of generality, we suppose the previous points to be in an alternating configuration of labels as we mentioned above. We can contemplate three possible scenarios for the added points:

i $p_{2k+1}, p_{2k+2} \in \ominus$

ii $p_{2k+1} \in \oplus, p_{2k+2} \in \ominus$, note²

iii $p_{2k+1}, p_{2k+2} \in \oplus$

¹ie, we can shatter $2k$ points but not $(2k) + 1$ points.

²Equivalent to $p_{2k+2} \in \oplus, p_{2k+1} \in \ominus$

case i

Since the previous $2k$ points could be shattered and there are no two contiguous \oplus labels in the previous set of $2k$ points, introducing two \ominus labels anywhere will not disrupt prior labeling if the intervals capturing the adjacent \oplus points are adjusted (shrunk) accordingly.

case ii

As above, the \ominus point will not disrupt prior labeling. The \oplus point will either fall beside another \oplus point where it can be included in the interval³ capturing the adjacent \oplus , at either end of the set, besides an \ominus point, or between the contiguous \ominus points that resulted from the insertion of $p_{2k+2} \in \ominus$; in either case, the $k + 1^{th}$ interval will correctly classify the newly inserted $p_{2k+1} \in \oplus$.

case iii

If the previous $2k$ points are labeled with alternating \ominus and \oplus , then one end of the set will have \ominus and the other \oplus . Thus on inserting points p_{2k+1} and p_{2k+2} one of them will necessarily fall beside another \oplus and, in the worst case, the other point could be placed at the end of the interval on the end with the \ominus , in which case the $k + 1^{th}$ interval would correctly classify it.

Thus $k + 1$ intervals shatter $2(k + 1)$ points. Conversely, with $2(k + 1) + 1$ points and the following configuration $\oplus, \ominus, \dots, \oplus$ we would not be able to shatter the set of points with $k + 1$ intervals.

Thus the inductive step holds.

Then, for this class with k intervals, $VC_{dim} = 2k$.

³Once the bounds of said interval have been adjusted

Question 2: KL Divergence

2.a $KL(P||Q) \geq 0, \forall P, Q$

Since $\log(x)$ is a concave function, in order to use Jensen's inequality as stated for convex functions, we make the expression convex by proving $-KL(P||Q) \leq 0, \forall P, Q$. We use $P(x_i) > 0, \forall P(x_i)$ for convenience.

$$\begin{aligned} -KL(P||Q) &= -\sum_{i=1}^m P(x_i) \log \left(\frac{P(x_i)}{Q(x_i)} \right) \\ &= \sum_{i=1}^m P(x_i) \log \left(\frac{Q(x_i)}{P(x_i)} \right) \end{aligned}$$

And by Jensen's inequality, taking $\frac{P(x_i)}{Q(x_i)}$ to be a random variable uniformly distributed over i , we can then write:

$$\begin{aligned} -KL(P||Q) &\leq \log \left(\sum_{i=1}^m P(x_i) \frac{Q(x_i)}{P(x_i)} \right) \\ &\leq \log \left(\sum_{i=1}^m Q(x_i) \right) \\ &\leq \log(1) \\ &\leq 0 \end{aligned}$$

2.b $KL(P||Q) = 0$?

When both distributions are equal, i.e. $P(x_i) = Q(x_i), \forall i$, $KL(P||Q)$ becomes:

$$\begin{aligned} KL(P||Q) &= -\sum_i P(x_i) \log \left(\frac{Q(x_i)}{P(x_i)} \right) \\ &= -\sum_i P(x_i) \log(1) \\ &= 0 \end{aligned}$$

Which makes sense since the divergence of two equal distributions should be zero.

2.c Max $KL(P||Q)$?

$$\begin{aligned} KL(P||Q) &= \sum_i P(x_i) \log \left(\frac{P(x_i)}{Q(x_i)} \right) \\ \lim_{Q(x_i) \rightarrow 0} (KL(P||Q)) &= \infty, \text{ for some } i, | P(x_i) \not\rightarrow 0 \end{aligned}$$

Which can be interpreted as the divergence between the true distribution $P(x)$ and the modeling distribution $Q(x)$ approaching infinite if $Q(x)$ cannot represent an event x_i with a non-zero probability in $P(x)$.

2.d $KL(P||Q) = KL(Q||P)$? **Justify**

No, by definition they are not symmetric. We show a counterexample to symmetry: suppose $x_i = \{0, 1\}$ with $P(0) = \frac{1}{3}, P(1) = \frac{2}{3}$ modelled by $Q(0) = \frac{1}{2}, Q(1) = \frac{1}{2}$. Then:

$$\begin{aligned}
 KL(P||Q) &= \frac{1}{3} \log \left(\frac{\frac{1}{3}}{\frac{1}{2}} \right) + \frac{2}{3} \log \left(\frac{\frac{2}{3}}{\frac{1}{2}} \right) \\
 &= \frac{1}{3} \log \left(\frac{2}{3} \right) + \frac{2}{3} \log \left(\frac{4}{3} \right) \\
 &= 0.056633 \\
 KL(Q||P) &= \frac{1}{2} \log \left(\frac{\frac{1}{2}}{\frac{1}{3}} \right) + \frac{1}{2} \log \left(\frac{\frac{1}{2}}{\frac{2}{3}} \right) \\
 &= \frac{1}{2} \log \left(\frac{3}{2} \right) + \frac{1}{2} \log \left(\frac{3}{4} \right) \\
 &= 0.058891 \\
 KL(P||Q) &\neq KL(Q||P)
 \end{aligned}$$

2.e **Prove** $KL(P(Y, X)||Q(Y, X)) = KL(P(X)||Q(X)) + KL(P(Y|X)||Q(Y|X))$

$$\begin{aligned}
 KL(P(Y, X)||Q(Y, X)) &= \sum_x \sum_y P(x, y) \log \left(\frac{P(x, y)}{Q(x, y)} \right) \\
 &= \sum_x \sum_y P(x, y) \log \left(\frac{P(y|x)P(x)}{Q(y|x)Q(x)} \right) \\
 &= \sum_x \sum_y P(x, y) \left(\log \left(\frac{P(x)}{Q(x)} \right) + \log \left(\frac{P(y|x)}{Q(y|x)} \right) \right) \\
 &= \sum_x \sum_y P(x, y) \log \left(\frac{P(x)}{Q(x)} \right) + \sum_x \sum_y P(x, y) \log \left(\frac{P(y|x)}{Q(y|x)} \right) \\
 &= \sum_x \sum_y P(x|y)P(y) \log \left(\frac{P(x)}{Q(x)} \right) + \sum_y P(y|x)P(x) \log \left(\frac{P(y|x)}{Q(y|x)} \right) \\
 &= \sum_x \sum_y \left[P(x|y)P(y) \log \left(\frac{P(x)}{Q(x)} \right) \right] + \sum_y \left[P(y|x)P(x) \log \left(\frac{P(y|x)}{Q(y|x)} \right) \right] \\
 &= \sum_x \log \left(\frac{P(x)}{Q(x)} \right) \sum_y [P(x|y)P(y)] + P(x) \sum_y \left[P(y|x) \log \left(\frac{P(y|x)}{Q(y|x)} \right) \right] \\
 &= \sum_x P(x) \log \left(\frac{P(x)}{Q(x)} \right) + P(x) \sum_y P(y|x) \log \left(\frac{P(y|x)}{Q(y|x)} \right) \\
 &= KL(P(X)||Q(X)) + KL(P(Y|X)||Q(Y|X))
 \end{aligned}$$

2.f Prove $\arg \min_{\theta} KL(\hat{P}||P) = \arg \max_{\theta} \sum_{i=1}^m \log P_{\theta}(x_i)$

First, we develop the *LHS* and we note that:

$$\begin{aligned} \arg \min_{\theta} KL(\hat{P}||P_{\theta}) &= \arg \min_{\theta} \sum_x \hat{P}(x) \log \left(\frac{\hat{P}(x)}{P_{\theta}(x)} \right) \\ &= \arg \min_{\theta} \sum_x \hat{P}(x) \log(\hat{P}(x)) - \sum_x \hat{P}(x) \log(P_{\theta}(x)) \end{aligned}$$

And since $\sum_x \hat{P}(x) \log(\hat{P}(x))$ depends solely on the observations and is fixed w.r.t. θ , we can then equivalently write:

$$\begin{aligned} \arg \min_{\theta} KL(\hat{P}||P_{\theta}) &= \arg \min_{\theta} - \sum_x \hat{P}(x) \log(P_{\theta}(x)) \\ &= \arg \max_{\theta} \sum_x \hat{P}(x) \log(P_{\theta}(x)) \end{aligned} \tag{1}$$

Then, for a given set of m observations $x_i, x \in X, i \in \{1, 2, \dots, m\}$ we can have $n \leq m$ unique values, which we index as $x_j, x \in X, j \in \{1, 2, \dots, n\}$. We highlight the difference between observation indexing, $i \in \{1, 2, \dots, m\}$ and value indexing, $j \in \{1, 2, \dots, n\}$. We can then substitute $\hat{P}(x_j) = \frac{|x_j|}{m}$ into equation 1:

$$\arg \min_{\theta} KL(\hat{P}||P_{\theta}) = \arg \max_{\theta} \sum_{j=1}^n \frac{|x_j|}{m} \log(P_{\theta}(x_j)) \tag{2}$$

Now we develop the *RHS*:

$$M.L.E(x_i, \theta) = \arg \max_{\theta} \sum_{i=1}^m \log(P_{\theta}(x_i))$$

Multiplying by a positive constant value does not change the result of \max_{θ}

$$\begin{aligned} &= \arg \max_{\theta} \sum_{i=1}^m \frac{1}{m} \log(P_{\theta}(x_i)) \\ &= \arg \max_{\theta} \sum_{j=1}^n \frac{|x_j|}{m} \log(P_{\theta}(x_j)) \end{aligned}$$

Which is precisely equal to equation 2. QED

Question 3: Implementation: K-means

K-means clustering has been implemented in the `A4_Q3.driver.m` matlab script. The file is included as part of this submission.

We reproduce, at the end of the answer, the main part of code for the reader's convenience, omitting most non-essential plotting and file manipulation instructions. The script and respective reprint includes some useful comments regarding the procedure.

The script outputs useful information both as simple text and in the form of plots. The text output is printed here and the plots can be seen in figures 1, 2 and 3.

The original and resulting images are also reprinted in figures 4a and 4b

There are 6 total clusters with pixels in them

Final cluster membership count is respectively:

4930 15190 52535 0 22075 0 40365 74917

The final centroids are:

241.22961	238.62515	233.86288
194.41159	136.33311	90.94365
136.26556	61.08973	10.10385
0.00000	255.00000	0.00000
157.29173	97.59398	51.43330
0.00000	0.00000	255.00000
78.92744	37.10829	13.07070
25.97800	23.23575	23.60599




Figure 1: The sum of all squared distances from all pixels towards their respective centroids over iterations.

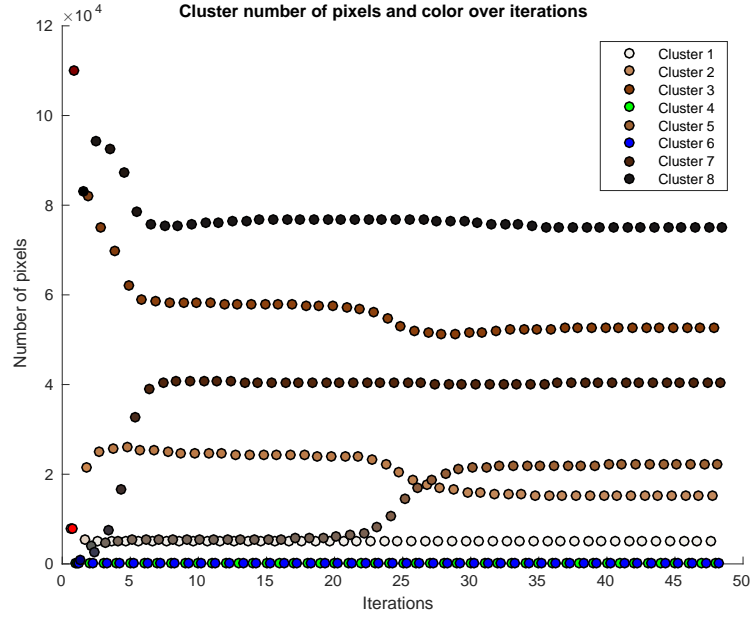


Figure 2: The number of pixels pertaining to each cluster per iteration. The marker colors represent the resulting color of centroids per iteration.

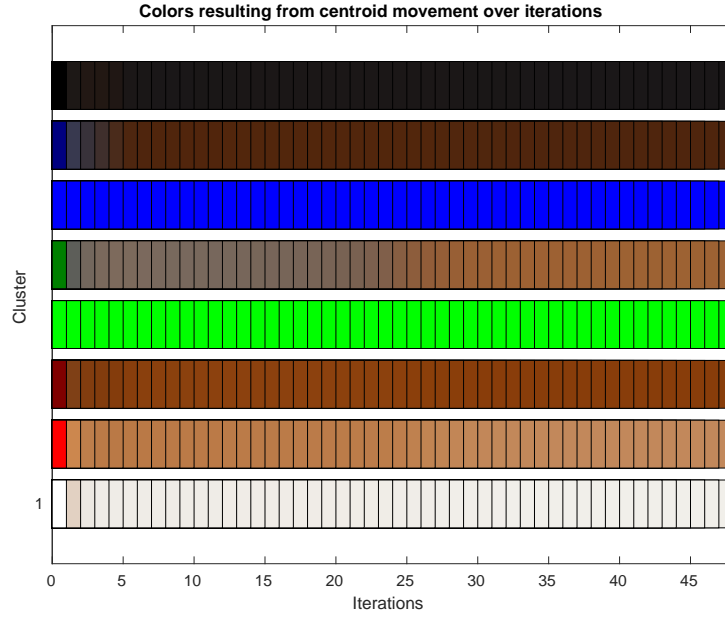


Figure 3: Centroids' resulting colors per iteration. Most centroids shift rapidly during the first few iterations, only cluster 5 seems to evolve more slowly, as does its membership count, see the bottom part of figure 2.



(a) original image - 256 colors



(b) image with 6 actual clusters for pixel color-value

Figure 4: The original image, left, and the clustered image to the right.

The k-means algorithm part of A4.Q3.driver.m:

```

24  init_matx = ...
25  [255, 255, 255; ...
26  255, 0, 0; ...
27  128, 0, 0; ...
28  0, 255, 0; ...
29  0, 128, 0; ...
30  0, 0, 255; ...
31  0, 0, 128; ...
32  0, 0, 0];
33
34  data = load('hw4-image.txt');
35
36  m = length(data(:,1));
37  d = length(data(1,:));
38  k = length(init_matx(:,1));
39
40  % guesstimated number of iterations for preallocation,
41  % This is mostly for efficiency as the loop will stop before if converged
42  ITERS = 50;
43
44  % Initial centroids or cluster means
45  k_means = repmat(reshape(init_matx', 1, d, k), m, 1, 1);
46  new_means = init_matx;
47  k_labels = 1:m;
48  k_settled = false;
49
50

```

```

51 % keep a trace of the important data through iterations
52 k_membership_counts = zeros(ITERS, k);
53 means_trajectory = zeros(k,d,ITERS);
54 sum_squared_dist = zeros(1, ITERS);
55 iteration_count = 1;
56
57 % keep k-copies of the data to quickly get the distance to centroids
58 unfolded_data = repmat(data, 1, 1, k);
59
60 % get the L2 norm of the row-slice of the difference between 8 pixel copies and centroids
61 slice_sq_norm = @(tensor)reshape(sum((tensor.^2),2), size(tensor, 1), size(tensor, 3));
62
63 % Repeat the process until done...
64 while (~k_settled)
65     loop = tic();
66     disp(sprintf('Clustering all pixels, iteration %d', iteration_count));
67     flush();
68     k_means = repmat(reshape(new_means', 1, d, k), m, 1, 1);
69     k_means_flat = reshape(k_means(1,:,:), d, k)';
70     [min_l2_sq, idxs] = min(slice_sq_norm(unfolded_data - k_means), [], 2);
71     sum_squared_dist(iteration_count) = sum(min_l2_sq);
72     k_labels = idxs;
73     disp(sprintf('Reassigning means'));
74     flush();
75     for kth = 1:k
76         temp = zeros(size(data));
77         kth_pixels_idx = (k_labels == kth);
78         temp(kth_pixels_idx, :) = data(kth_pixels_idx, :);
79         temp(~kth_pixels_idx, :) = [];
80         k_membership_counts(iteration_count, kth) = size(temp, 1);
81         if (size(temp, 1) >= 1)
82             new_means(kth, :) = mean(temp);
83         end
84     end
85     means_trajectory(:, :, iteration_count) = k_means_flat;
86     disp(sprintf(...
87         'Norm of the difference between this iteration and last''s means: %d', ...
88         norm(new_means-k_means_flat)))
89     flush();
90     disp(sprintf('iteration %d took %d seconds.', iteration_count, toc(loop)));
91     flush();
92
93 % See if we're done
94 if (sum(sum(new_means ~= k_means_flat)) == 0)
95     k_settled = true;
96 else
97     iteration_count = iteration_count + 1;
98 end
99 end
100
101 % Now convert pixels to their respective centroid
102 clustered = data;
103
104 for kth = 1:k
105     clustered(k_labels == kth, :) = k_means(k_labels == kth, :, kth);
106 end

```

Question 4: K-medoids - advantages and disadvantages vs K-means

K-medoids, which minimizes the dissimilarities inside the cluster⁴, has two main advantages over K-means which minimizes the sum of squared distances from points to means in each cluster. These advantages are:

- (i) it can use any measure of similarity between points, making it more flexible than K-means which can only use euclidean distance between vectors
- (ii) K-medoids partitioning is much more robust to outliers than K-means given the use of medoids as opposed to means

Conversely, the main disadvantage of K-medoids against K-means is that, for k, n, i clusters, iterations and data points respectively, K-medoids' runtime cost is $O(k \cdot i \cdot n^2)$ whereas K-means' is $O(k \cdot i \cdot n)$.

To put this in context: where a simple example such as the one from question 3 has a K-means time complexity bounded by $8 \times 48 \times 407 \times 516 \approx 352 \times 10^9$ operations, if we assume convergence also at 48 iterations for K-medoids, its time complexity would be in the order of $8 \times 48 \times (407 \times 516)^2 \approx 75 \times 10^{15}$ operations! Making K-medoids a pertinent choice only when applications are not time-sensitive where the added benefits of the use of medoids are desired.

⁴however these dissimilarities might be defined in a given context.